Faster Compression of Patterns to Rectangle Rule Lists

Ian Albuquerque Raymundo Da Silva[1], Gruia Calinescu[2], and Nathan De Graaf[3]

[1] Pontifical Catholic University of Rio de Janeiro, Brazil
`ian.albuquerque.silva@gmail.com`
[2] Illinois Institute of Technology, Chicago, Illinois
`calinescu@iit.edu`
[3] Iowa State University, Ames, Iowa
`nathandegraaf@gmail.com`

Abstract. Access Control Lists (ACLs) are an essential security component in network routers. ACLs can be geometrically modeled as a two dimensional black and white grid; our interest is in the most efficient way to represent such a grid. The more general problem is that of Rectangle Rule Lists (RRLs), which is finding the least number of rectangles needed to generate a given pattern. The scope of this paper focuses on a restricted version of RRLs in which only rectangles that span the length or width of the grid are considered. "Applegate et al."'s paper Compressing Rectilinear Pictures and Minimizing Access Control Lists gives an algorithm for finding the optimal solutions for strip-rule RRLs in $O(n^3)$ time. Following the structure of "Applegate et al."'s algorith, we simplify the solution, remove redundancies in data structures, and exploit overlapping sub-problems in order to achieve an optimal solution for strip-rule RRLs in $O(n^2 \log n)$ time.

1 Introduction

We consider the following problem of generating patterns by drawing rectangles. A target pattern is given as a grid of black and white cells. We begin with a white grid and place solid black or white rectangles, each rectangle covering any previously placed rectangles. Placing arbitrary sized rectangles makes the problem NP-Hard, so we consider only placing rectangles that extend either the full height or width of the grid. Our goal is to find the smallest number of rectangles required to create the target pattern.

1.1 Problem Definition

Define a pattern P to be an n_R by n_C grid of black and white squares. Let $n = n_R + n_C$.

A rectangle strip-rule on P is either a black or white rectangle that extends from one side of the pattern to the opposite side. Precisely a rectangle in P is either a set of contiguous rows of P or a set of contiguous columns of P

A rectangle strip-rule list (a SRRL) that generates P is an ordered list of rectangle strip-rules, that when applied in ordered to a blank (white) grid the size of P create the target pattern. See Fig. **??** for an example of a pattern generated by a SRRL.


Fig. 1. A pattern generated by a SRRL of 3 elements

We say a pattern P is a strip-rule pattern if there is a SRRL that generates P. Note that not every pattern P is a strip-rule pattern. See Fig. **??** for an example of a pattern that is not a strip-rule pattern.

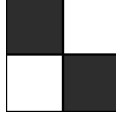An SRRL is considered optimal if it has the minimum number of rectangle strip-rules of any SRRL that generates P.

Fig. 2. The 2×2 checkerboard: an example of a pattern that is not a strip-rule pattern

1.2 Problem Background

Rectilinear Pictures and Access Control Lists. The method of stacking rectangles to create patterns has applications in both graphics and network routers. A common method for drawing graphics is to allow the user to repeatedly apply a rectangle tool to a blank canvas (See Xfig or PowerPoint). Each rectangle is of a solid color and covers everything in a defined rectangular region. This sequence of rectangles drawn can be represented by a rectangle rule list (RRL). The problem of finding the minimum length RRL needed to create a given pattern is a generalization of the SRRL problem we explore.

Alternatively, instead of being given an n_R by n_C grid as input, one can start with the numbers n_R and n_C, and a list of m rules. The goal is to find a minimum-length list that gives the same pattern as the input list. As shown in an extended version of [**ACJKLW07**], it is possible to construct the n_R by n_C grid in time $O(n_R \cdot n_C + m^2)$.

One important application of RRLs is in access routers. An Internet Service Provider might use access control lists (ACLs) on network router line cards in order to choose whether to forward or drop a packet based on the sending or receiving IP address. This decision could be answered by checking a two-dimensional Boolean array based on the sender and receiver IP addresses. This problem can be translated into a restricted version of our RRL problem. An in depth discussion of this translation and its properties is discussed in detail in "Applegate et al.""s paper [**ACJKLW07**]. In ACLs minimization, the rectangles are of the certain form: n_R and n_C are powers of two, rows and columns are indexed with binary strings from 0^w to 1^w, and each rectangle must be given by two prefixes y and z of length at most w each; the actual rectangle consists of the rows whose index has y as a prefix and the columns who have z as prefix.

Related Work. Unfortunately, the general problem of finding minimum length RRLs has been shown to be NP-Hard by [**ACJKLW07**]. Instead, we work on a restricted version of the problem in which any rectangle rules applied must extend either the height or width of the original pattern and only black or white rectangles are allowed. This 2-color strip-rule problem was originally posed in "Applegate et al.""s paper [**ACJKLW07**], where an optimal solution is given in $O(n^3)$-time, where n is the sum of the height and width of the original pattern. [**ACJKLW07**] obtains an optimal solution to the strip-rule version of the ACL problem with a similar (but more complicated) $O(wn^3)$-time algorithm. A 1.5 ratio approximation algorithm is given for the problem in $O(n^2)$-time. While there exist numerous results related to various other restricted problems regarding RRLs, we know of no other work related to this 2-color strip-rule problem.

ACL can be used in firewalls [**KRRW12**]. [**KRRW12b**] introduce axioms with the goal of creating and analyzing algorithms for optimizing the rewriting of rules in Software Defined Networks. [**LMT10**] consider classifiers in dimensions higher than two, which reduce to either ACL or RRL lists in dimension two. They propose and experimentally evaluate heuristics without performance guarantees, as well as relate these classifiers to the Firewall Decision Diagrams of [**GL07**]. [**PL14**], and [**CDG16**] also consider higher dimensional classification based on rules. [**KLRW13**] and [**ZNHBM14**] introduce more general rule-minimization problems. [**SK10**] uses rule minimization as a start for solution to an extended problem. Efficiently removing redundant rules has been proposed in [**SK10b**].

The complexity of finding minimum length ACLs in two dimensions is still unknown, but with an arbitrary number of dimensions, this problem is NP-hard [**KNCER13**]. [**ACJKLW07**] gives a $O(\min(m^{1/3}, 1^{/2}))$-approximation algorithm for finding minimum length RRLs, where m is the length of the input *RRL*. This is still the best published ratio.

[**DLT16**] provides heuristics for higher-dimensional ACL and RRL minimization ([**DLT16**] calls the ACLs minimization prefix-ACL, while range-ACL is their terminology for finding minimum length RRLs), on the way improving the approximation ratios provided by [**ACJKLW07**] for ACLs minimization. The approximation algorithms of [**DLT16**] and [**ACJKLW07**] use as a subroutine the strip-rule version that we study.

Our Results. Using the structure defined in "Applegate et al.""s [**ACJKLW07**] $O(n^3)$ exact algorithm for 2-color SRRLs, we give an improved $O(n^2 \log n)$ exact algorithm for the SRRL problem. Given the similar structures of

the RRL and ACL optimal solutions given by [**ACJKLW07**], we expect our solution can be extended to improve the running time of an exact algorithm for the strip-rule ACL problem from $O(wn^3)$ to $O(wn^2 \log n)$. As strip-rule ACLs occur in high percentage of ACL minimization cases [**ACJKLW07**], this is one reason to study SRRL. Another two reasons are: we believe SRRL is a natural problem, and SRRL is used in the approximation algorithms for RRL minimization.

Our result is obtained by digging deeper in the structure of the dynamic programming of [**ACJKLW07**] combined with use of geometric data structures to speed up the process. We use fast two-dimensional orthogonal range queries, using existing data structures (a time bound of $O(\log n)$ time per query being textbook material [**KOS2000**]).

1.3   Organization

We begin by defining all structures and concepts relevant to finding the optimal SRRL. Next we give a detailed explanation of our $O(n^2 \log n)$ optimal algorithm. Finally, we state all theorems that make the algorithm possible and analyze the algorithm's time and space complexity.

2   Preliminaries

We begin by exploring the strategies and tools for finding an optimal SRRL. The Pick-Up-Sticks algorithm will detail the basic structure used to find an SRRL, but not necessarily the optimal SRRL. Then we will explain the concept of equivalence classes, which can be grouped into segments. All of these concepts will be put together in our algorithm for finding the optimal solution.

2.1   The Maximum Pick-Up-Sticks (MPUS) Algorithm

As proposed by "Applegate et al.""s paper, we can build an algorithm for finding whether a pattern is a strip-rule pattern and if so, finding an SRRL that generates it.

The Pick-Up-Sticks (PUS) Algorithm. The Pick-Up-Sticks algorithm of [**ACJKLW07**] is an algorithm for generating a SRRL of a pattern P if such a list exists. The idea is to pick up monochromatic rows and columns in order to build the SRRL backwards. Every time we pick a row or column, we color it gray.

Let P be a black and white pattern. We define a column or row in P as being pseudo-monochromatic if it is composed of only gray and white cells or black and gray cells. Note that monochromatic columns and rows are also pseudo-monochromatic.

The Pick-Up-Sticks algorithm builds an SRRL as follows:

While there are still black cells, choose a pseudo-monochromatic column or row and color all its cells gray. After a row or column is colored gray, add a rectangle that corresponds to covering that row or column with whatever non-gray color was left in that row or column to the beginning of our SRRL. Note that, in our problem definition, we begin applying SRRLs with a white grid, so we can stop picking up sticks when no black cells exist. If, on the other hand, we modify the problem to begin with a grid of some third color, we would only stop picking up sticks when all cells are gray. There may be a difference of one rule between the optimum solutions to these two problems (white grid or some third color grid) with the same target pattern. For the sake of symmetry, from now on we use this *modified version* of the problem. The method works with minor modifications for the original version.

The algorithm may not succeed, as it may not find a pseudo-monochromatic column or row. If the algorithm does end in a all-gray grid, the algorithm has "picked up" the rectangles of the SRRL in reverse order. Figure **??** gives a representation of one possible execution of the Pick-Up-Sticks algorithm and Fig. **??** shows the generation of the original pattern from the resultant SRRL.

The Pick-Up-Sticks algorithm is guaranteed to generate an SRRL if one exists [**ACJKLW07**] (included in Appendix, Subsection **??**), but it is not necessarily the optimal SRRL of minimum length. At any stage of the algorithm there will be more than one option on which row or column to pick up. Different choices can lead to different sizes of SRRLs. In the next sections we will discuss on how "Applegate et al.""s paper narrows down the number of choices for each stage in order to find the minimal SRRL.
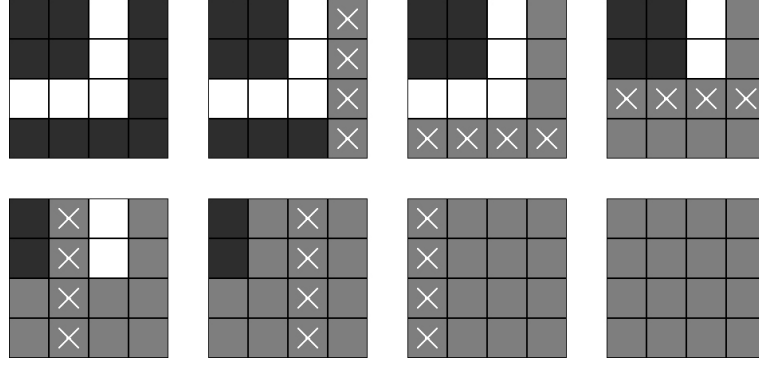
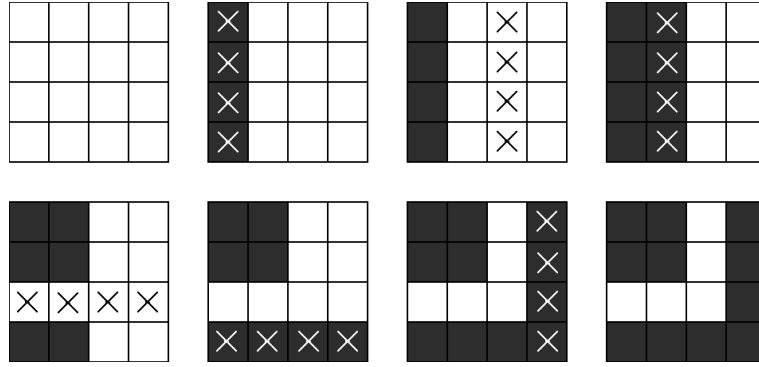**Fig. 3.** The execution of the Pick-Up-Sticks algorithm on a strip-rule pattern.

Fig. 4. The SRRL generated by the execution of the MPUS algorithm on the pattern of the Figure **??**

Improving by Picking Maximal Sticks. One important observation by "Applegate et al."'s paper on the Pick-Up-Sticks (PUS) algorithm is that there can be no harm in always using maximal strip-rules. A maximal strip-rule is a strip-rule that picks up a maximal contiguous sequences of pseudo-monochromatic rows or columns. Following [**ACJKLW07**], we call such a contiguous sequence a *block*.

This is possible because replacing a non-maximal strip-rule by the maximal one that contains it does not affect the pseudo-monochromaticity of any later rule.

Hence, we may always use the Maximal Pick-Up-Sticks (MPUS) algorithm instead of the Pick-Up-Sticks (PUS) algorithm. In this new algorithm, once we choose the pseudo-monochromatic column to pick, we pick the maximal contiguous set of pseudo-monochromatic columns that contains it.

2.2 Equivalence Classes

We now introduce the concept of equivalence classes, as proposed by "Applegate et al."'s paper. This definition will allow us to give structure on the order of which columns and rows should be picked-up during an execution of MPUS in order to find an optimal SRRL.

Definition and Properties. Given a pattern P, we define two rows or columns as being of the same equivalence class if and only if they are both of the same type (row or column) and their cells have the exact same colors in P, in the same order; that is, they are identical.

We will denote a column equivalence class in a pattern P as being C_x where x is the number of black cells in the original pattern. Analogously, we will denote R_y as being the row class with y black cells in the original pattern. See Fig. **??** for an example of a pattern and its equivalence classes. It follows from the monotocity property (Theorem **??** in Subsection **??** of the Appendix) that all the columns in C_x are identical in all positions, for all x, and the same property holds for all the rows of R_y, for all y.

Note that every row and column belongs to exactly one equivalence class. Also, if two columns belong to a equivalence class in the beginning of the MPUS algorithm, they will remain of the same class until one of them

Fig. 5. Equivalence classes of a black and white pattern P.

is picked up by the algorithm. This is since picking up other columns does not change these two columns at all, while picking up rows modify these two columns in exactly the same way. The same property holds for two rows that belong to a equivalence class.

We also define an equivalence class as being *active* during the execution of the MPUS algorithm if some member of that class is pseudo-monochromatic but not all gray.

Active Classes Properties. "Applegate et al."'s paper shows that during the execution of the MPUS algorithm on a black and white strip-rule pattern, there are always exactly two active classes at any given time. The intuition on why this happens is that in order for a new class to become active, all the members of an old active class need to be picked up. A formal proof of this property appears later in Proposition **??** in the Appendix (Subsection **??**); this fact is also used in [**ACJKLW07**]. Also, the same proposition shows that the two active classes are either a row and a column class of the same color or both classes of the same kind (rows or columns), being one of each color.

Embedded Rows and Columns. We can start improving the MPUS algorithm by introducing the concept described by "Applegate et al."'s paper as embedded rows and columns. This will allow us to pick up more rows and columns using the same number of rectangles.

First, we will define the concept of a embedded column or row at a given stage of the MPUS algorithm. Given a point during the execution of the MPUS algorithm, let b be a column of an active equivalence class B. Let a_1 be the first column to the left of b that does not belong to B and has not yet been picked up. Similarly, let a_2 be the first column to the right of b that does not belong to B and has not yet been picked up. We say that b is *embedded* in equivalence class A if both a_1 and a_2 belong to that class A. The definition is analogous for rows. Note that columns can be embedded only in column classes and rows only in row classes. See the three leftmost columns of Fig. **??** for a simple example of a white column embedded in two black columns.

Fig. 6. Example of a pattern P where some members of C_0 are embedded in some members of C_2.

If during the execution of the MPUS algorithm we have the option to pick up two column-blocks that embeds a set of columns of another class, we could pick up the two column-blocks that embeds the third one with two rectangles and then the embedded one, totaling three rectangles. However, it is better to pick the embedded block first with one rectangle, and then pick the two blocks that were previously embedding with only one rectangle, thus using a total of two rectangles for the same set of columns. As an example, in Fig. **??**, one cannot benefit by picking up the first and the third column, when one can pick up first the second column, followed by picking up the block of the first three columns.

Equivalence Classes Ordering. As we pick up all the members of a column class, we change the rows of the pattern. Since the other columns remain intact, the new class to become active must be a row class. By using the same argument for row classes, we see that whenever we pick up all the members of a class, the next class to become active is of the opposite type. This is formally proven in Proposition **??** in the Appendix. Also (same proposition), if we pick up all the members of an active class that was pseudo-monochromatic on a given color (for instance, black), the next class to become active needs to be of the opposite color (for instance, white). This happens because if it was of the same color, the other class would already have been active.

We can order the equivalence classes of both rows and columns in a given pattern P by the number of black cells on it. In that ordering, we can see (Proposition **??** in the Appendix) that the all the rows that two consecutive column classes differ form exactly an equivalence class. This holds analogously for columns. This implies (Proposition **??** in the Appendix) that the numbers of columns and rows classes differ by at most one. Using that information, we can set up the following ordering for the equivalence classes (see also Proposition **??** in the Appendix):

1. Start with the column class with only white cells. If such class does not exist, start with the row class with only black cells (the existence of this row class follows immediately from the monotonicity property - Theorem **??** in the Appendix).
2. Alternate columns and rows, putting the columns in ascending order of black cells and the rows in descending order of black cells.

Let $C_{x_0}, C_{x_1}, \cdots, C_{x_{(N_C-1)}}, C_{x_{N_C}}$ be the ascending ordering of column classes by the number of black cells. Let $R_{y_0}, R_{y_1}, \cdots, R_{y_{(N_R-1)}}, R_{y_{N_R}}$ be the ascending ordering of row classes by the number of black cells. Using the rules above, we should get an ordering similar to $C_{x_0}, R_{y_{N_R}}, C_{x_1}, R_{y_{(N_R-1)}}, \cdots, R_{y_1}, C_{x_{(N_C-1)}}, R_{y_0}, C_{x_{N_C}}$. We will call this ordering the *hierarchical array* of the pattern P. We will also refer to this ordering of equivalence classes as $E_1, E_2, \cdots, E_N$. Note that N_C, the number of column classes, is at most n_C, the number of columns, and the same property holds for rows.

One important property of the hierarchical array of a pattern is that if E_x and E_y, $x < y$, are the active classes during the MPUS execution, all classes that come before E_x in the sequence or after E_y must have already been picked up (otherwise, E_x and E_y would not be active). This is proven in Proposition **??** in the Appendix. Also from **??**, if a class is between those two classes, none of the rows/columns of the in-between class are pseudo-monochromatic. The hierarchical array is a useful extension of Observation 5 on page 1070 of [**ACJKLW07**].

2.3 Segments

During MPUS, there are always two active classes. A "segment" succinctly represents the blocks that are available for pick up while a certain pair of classes are active. These segments can be translated into nodes of a graph, that can be used for finding the optimal solution, as described below. Note that our representation of segments is a slightly condensed version of the one described in the original "Applegate et al."'s paper [**ACJKLW07**]. The difference is that the original definition is a five tuple with two redundant fields. See Subsection **??** of the appendix for a proof that these fields are redundant. It is not here where the running time is improved.

Definition and Properties. Whenever a new class becomes active during the execution of the MPUS algorithm on a pattern P, we will define a *segment S* as being a tuple of three elements (E_x, E_y, U) where E_x and E_y are the active equivalence classes and U is the subset of members (rows or columns) left unpicked of E_y ($U \subset E_y$).

We can represent the sequence of actions of the MPUS execution as a sequence of segments, as described in this paragraph. Given the pick-up order obtained through the execution of the MPUS, let's create a segment (E_0, E_N, E_N) for the starting state of pattern P and a segment (E_x, E_y, U), $U \subset E_y$, for every time a new class E_x becomes active; we say that the SRRL *includes* the segment whenever such a segment becomes active during the MPUS execution corresponding to the SRRL.

Since E_x is the new class, all of the members of E_x are still in the pattern. E_y, in the other hand, is a class that was already active. It may have some of its members already picked-up. We store that information in U by maintaining the members that have not been picked up yet. Note also that any classes that are between E_x and E_y in the hierarchical array of P have not become active yet, resulting in the fact that all of their members are still there. At the same time, members that come before and after the interval delimited by E_x and E_y have all been picked up. Hence, for each one of those segments that are created whenever a class becomes active it is possible to reconstruct the pattern at that moment of the execution of the MPUS algorithm.

Segment Branching. "Applegate et al.'"s paper [**ACJKLW07**] shows that if we are looking for an optimal solution we can narrow down the number of segments significantly by reducing the number of options a segment can branch to two.

Suppose a minimum-length SRRL for a strip-rule pattern P includes the segment $S = (E_x, E_y, U)$ and that to get to the next segment, if it exists, we pick up all members of E_x (possibly picking up some members of $U \subset E_y$) while at least one member of U remains active.

- If $U \subset E_y$ has a member that is not embedded in E_x, there exists a minimum-length SRRL that includes the segment $S = (E_x, E_y, U)$ and that to get to the next segment you pick up all the members of $U \subset E_y$ that are embedded in E_x, followed by picking all members of E_x.
- If all members of $U \subset E_y$ are embedded in E_x, there exists another minimum-length SRRL that includes the segment $S = (E_x, E_y, U)$ and that to get to the next segment you pick up all the members of $U \subset E_y$ (thus, finishing picking up every member of E_y in the pattern). Note that in this case, one can get to the next segment by picking up all the members of $U \subseteq E_y$ first.

The same property holds with the roles of E_x and U switched. In this case, we assume the next segment from S, if it exists is originally reached by picking up all members of $U \subset E_y$.

This happens because we can always pick embedded members at no extra cost, if one is to pick all the blocks of the embedding class, as shown in Proposition **??** of the Appendix. Moreover, suppose we have two sets U_x and U_y of active blocks of the two active classes E_x and E_y. To finish picking up blocks (except at the very end), we must reach the situation that either E_x or E_y are not active anymore; say E_x becomes not active, while E_y is still active. Picking up any block from the class E_y that is not embedded in E_x can be done, without loss of optimality, later. So we can assume that only the members of U_y that are embedded in E_x are picked up while E_x is active.

Note also that for every segment we could pick the latest class that became active or the oldest, a total of two options. We can then say that every segment will branch into two other segments. On both cases, we can determine the next segment reached by eliminating every embedded member of the class not being picked-up, followed by eliminating every member of the class being picked-up. The only exception to this, is if when we pick every embedded member of the class not being picked-up we end up picking all of its members, thus already reaching a new segment. In that case, the original segment branches into only one segment, instead of two.

Note that there could be a segment that does not have any next segment. In this case, the segment is a segment at the very end of minimum-length SRRL. If we pick up any of either class, we will obtain an all-gray grid and the algorithm will be over. Also, note that for this to happen we must have that E_x and E_y are next to each other in the hierarchy array, and picking-up one of them will also result in picking-up the other one.

Segment Graph. We build a graph using the segments as nodes, as described by the following structure proposed by [**ACJKLW07**]:

- Start with the segment (E_0, E_N, E_N), that represents the grid as it is in very beginning of the MPUS algorithm. This will be the starting node.
- For every node (segment) in the graph, generate the two segments (possibly one, as described in the previous subsection) that it branches into and add a directed edge from it to the new generated nodes. There will be one segment for each of the two options of picking up the newest or the oldest of the two classes in the segment. On both cases, we can determine the next segment reached by eliminating every embedded member of the class not being picked-up, followed by eliminating every member of the class being picked-up.
- Define the cost for each edge as the number of rules in the MPUS algorithm used to go from one segment to another.
- Create an artificial node corresponds to an all-gray grid. This will be the end node of the graph. Every time a segment has its two classes adjacent to each other in the hierarchy array, create an edge from it to this all-gray grid node. This edge should have cost one (since we solve the modified version); if we were to solve the original version, this edge should have cost zero if the members of the classes of the segment were white and one otherwise.

Because of what was discussed in the Segments Branching section, at least one of the paths from the starting node to the end node of the graph will be a minimum-length SRRL. Note that having (E_N, E_0, E_0) as the starting node will lead to a graph that represents the same possible steps for the MPUS because it also represents the same starting pattern. Note also that some nodes may be reached by more than one node. Every node will only branch

into nodes that correspond to patterns with fewer total of columns and rows. This implies that there are no cycles in this graph.

We say that one segment reaches another if there is a path in this graph from that segment to the other one. The distance from one segment to another is the sum of the costs of the edges that compose the path between them, if such exists.

One very important property of this graph that will be proven later is that if two segments (E_x, E_y, U) and (E_x, E_y, U') are reachable from the the starting node of the graph, then either $U \subset U'$ or $U' \subset U$. This property will be referred as the Containment Lemma and appears in Subsection **??** in the Appendix. We will see that this implies that the number of nodes in this graph is $O(n^2)$ where n is the number of rows and columns in the original pattern.

3   Algorithm

The flow of our algorithm closely follows that of the $O(n^3)$ algorithm given by "Applegate et al.". We similarly create a segment graph of the reachable states in a MPUS of the pattern. However, we create this graph faster by grouping similar segments into S-groups, which can be processed together.

3.1   Setup

We begin by reading the input pattern into an equivalence class list and creating the first node in our segment graph.

Equivalence Class List. We are given as input an $n_R \times n_C$ grid of black and white cells. Assign each of the rows an index from 1 to n_R from top to bottom, similarly each column, an index from 1 to n_C from left to right. Group rows and columns into equivalence classes and order the classes by number of black cells as previously described in the Subsection **??**. Precisely, construct the hierarchical array:

$$C_{x_0}, R_{y_{N_R}}, C_{x_1}, R_{y_{(N_R-1)}}, \cdots, R_{y_1}, C_{x_{(N_C-1)}}, R_{y_0}, C_{x_{N_C}}$$

(this array may start and/or end with row equivalence classes, instead of columns as above).

We will need to determine from a range of column indexes which columns are also in a range of equivalence classes, and the similar question with rows instead of columns. A data structure such as an orthogonal range tree [**KOS2000**] accomplishes these queries in $O(\log n)$ time with $O(n \log n)$ space and setup time.

S-group. Define an S-group (E_1, E_2) (for equivalence classes E_1 and E_2, where the order of E_1 and E_2 in the tuple above matters) to be a collection of all segments $\{C_1, C_2, U'\}$ such that $C_1 = E_1$ and $C_2 = E_2$. To completely represent each of these segments, an S-group (E_1, E_2) maintains a list of segments S and a master list U. This list U will be used to keep track of each of the segment's individual list u, as described below. This allows each segment to be stored in constant space, while the S-group gets stored in $O(|E_1| + |E_2|)$ space.

The list U contains the index of every member of E_2 ordered such that for every segment $\{E_1, E_2, U'\}$ in the S-group, the first $|U'|$ members of U comprise the same set as U'. Such an ordering is guaranteed to exist by the Containment Lemma from Subsection **??**. We generate this ordering by noting that indexes appearing in more segments appear earlier in the list U.

S, the list of segments maintained by the S-group, keeps these segments ordered by the size of their set U'. Each segment s_i will have three values, $u(s_i)$, $d(s_i)$, and $p(s_i)$. $u = |U'|$. d gives the minimum number of sticks required to reach the state of segment starting from the original pattern. We want to maintain the shortest path to each node, so we use the value p to store a reference to the segment right before this one on this path.

Origin S-group. We begin with the S-group (E_1, E_2) with $U = \overline{E_2}$ and S comprised of a single segment $(u, d, p) = (|\overline{E_2}|, 0, null)$, where $\overline{E_2}$ represents the complete set of the members of E_2. The classes E_1 and E_2 are the first and last members of the hierarchical array. We will explore outwards from this S-group in a breadth first fashion, enumerating all reachable segments. The graph on the S-groups is guaranteed to be acyclic because the number of picked up classes is strictly increasing along each edge.

3.2 Finding the Next S-group

Suppose we have some S-group (E_1, E_2) with associated lists U and S. Without loss of generality we will assume E_1 to be a column class. We wish to generate all S-groups reachable from this S-group. The two S-groups immediately reachable from this S-group correspond to completely removing either class E_1 or E_2.

To determine the cost of removing a class E_i from a segment s_j, we count how many members of E_{3-i} can be picked up for free, as described in detail later. If two members of E_i may be picked up with the same stick while E_1 and E_2 are the active classes, we will call the two members *contiguous*. It follows from Proposition **??** that a pair of columns is contiguous in an S-group (E_1, E_2) if every column between them is either already picked up or of type E_1 or E_2.

Before the class E_i can be removed from a segment (E_1, E_2, U') (all the members of the class being picked up), we pick up every embedded member of E_{3-i}. To determine the cost of moving to the next segment, we must count how many of the embedded columns are also contiguous. Then any embedded columns must be removed from either the set U' (if E_1 is removed) or E_1 (if U' is removed) in the next segment and, by doing this for all segments, in the next S-group. In fact, we process one S-group at a time, and obtain segments of another at most two S-groups, as explained below.

Because E_1 and E_2 are not symmetric (E_2 comes with its set U), we will describe the process of removing classes E_1 and E_2 separately.

Remove E_1. Let E_3 be the equivalence class adjacent to E_1 in the equivalence class list, which has not yet been removed. The S-group corresponding to the result of removing E_1 will be (E_3, E_2) with lists $\hat{U}$ and $\hat{S}$.

Count Contiguous Consecutive Pairs of Columns of E_1. To determine the distances to the segments of the next S-group, we must count the number of contiguous E_1 ranges.

To do so, we must determine the range of equivalence classes which have not yet been picked up. Let E_1 be the column class C_a. If E_2 is also a column class call it C_b. If E_2 is a row class, then C_b is the column class adjacent to E_2 in the equivalence class list that has already been picked up. We assume $a < b$, with the other case being symmetric. By Proposition **??**, a column class C_i has not been picked up iff $i \in (a, b)$.

For each pair of adjacent members of E_1 we check to see if they are contiguous. Precisely, for each member $e_i \in E_1$, we query our column range tree for members in the range $[a+1, b-1]$ with index in the range $[e_i, e_{i+1}]$, where e_{i+1} is the column of E_i that has the smallest index among those with index higher than e_i (e_{i+1} is the next column of E_i after e_i). Let c be the number of pairs of consecutive members of E_1 which are also contiguous. The value $(|E_1| - c)$ corresponds to the number of sticks required to completely remove E_1 after any embedded columns of E_2 are picked up.

Get Embedded E_2 Blocks. If E_2 is also a column class, we must count and remove any members of E_2 embedded in members of E_1. Similar to the way we checked that columns of E_1 were contiguous, we will check if adjacent columns of E_1 are both contiguous and contain columns of E_2. In order to accurately keep track of the cost to remove these embedded columns we must give each embedded column a tag based on which two columns surrounded the embedded column. Two members of E_2 can be picked up with the same stick if and only if they have the same tag.

When an embedded column is found, tag it with the index of the E_1 column to its left. Then add this column to a list, B, sorted by the column's index. After we have found all of the embedded columns, we are ready to generate $\hat{U}$ and $\hat{S}$.

Build the Next S-group. Begin with $\hat{U}$ as an empty list. We also need a set T to keep track of which tags have been accounted for. Iterate through U, searching for each member m of U in the set of embedded columns, B. If $m \notin B$, append m to the end of $\hat{U}$. Otherwise add m's tag value to the set T if it is not already there. Once we have checked all of the members U' in a segment s_i (U' being the first $u(s_i)$ elements of U), add a new segment $\hat{s}_i$ to $\hat{S}$ with the following values.

$$u(\hat{s}_i) = |\hat{U}|$$

$$d(\hat{s}_i) = d(s_i) + |T| + |E_1| - c$$

$$p(\hat{s}_i) = s_i$$

(note that $|\hat{U}|$ and $|T|$ are computed for the sets $\hat{U}$ and T exactly when finish processing the last element of U' while iterating through U; $\hat{U}$ and T can change later on). If $\hat{S}$ already contains a segment $\hat{s}_j$ with $u(\hat{s}_j) = u(\hat{s}_i)$,

then keep only the segment with shorter distance d in the set $\hat{S}$, and remove the other one. As an aside, one can see that if we don't remove duplicates, if two segments from the same S-group s_i and s_j have $u(s_i) > u(s_j)$, then $u(\hat{s}_i) > u(\hat{s}_j)$, which is used in the proof of the Containment Lemma (Subsection **??** in the Appendix).

Remove E_2. Let E_3 be the equivalence class adjacent to E_2 in the equivalence class list, which has not yet been removed. The S-group corresponding to the result of removing E_2 will be (E_3, E_1) with lists $\hat{U}$ (whose elements are members of the class E_1) and $\hat{S}$.

Count Contiguous Consecutive Pairs of Columns of E_2. To count contiguous ranges of E_2, we must count the contiguous ranges within each segment separately. We use a counter c, initially set to 0. Fortunately, if a pair of columns is contiguous in one segment, it is also contiguous in all larger (with bigger value of u) segments. For each index $u_j \in U$ (elements of E are indices of columns), insert u_j into a sorted list of indexes, I. Get the predecessor and successor of u_j in I and check if these ranges from u_j to its neighbors are contiguous. If either of these ranges exist and are contiguous then this column can be picked up for free. If not, we increment our cost counter c. Once we have iterated over the first $u(s_i)$ blocks, we save the state of our counter in a value $c_i = c$. This value corresponds to the cost to completely pick up the columns in the segment s_i after any embedded members of E_1 have been picked up. Proceed (with c pussibly increasing) until we finish the list U.

Get Embedded E_1 Blocks. If E_2 is also a column class, we must count and remove the embedded columns of E_1. Each segment could have a unique number of embedded columns, where the larger the segment, the more columns can be embedded and the smaller the resulting segment will be in the next S-group. As an aside, this is an argument used in the proof of the Containment Lemma from **??**. To generate the ordering of $\hat{U}$ and the values in $\hat{S}$ and we must carefully count the embedded columns.

We iterate through the columns of U keeping track of which columns are embedded and will get picked up. To help us with this task, we start with a sorted list V of the indexes of E_1, an empty list for $\hat{U}$, and a counter d, initialized to $d = 0$, to measure the number of required sticks. Also start with B, a set of columns of E_1, initialized as the empty set.

Similar to the way we counted contiguous blocks, for each column $u_i \in U$ we insert u_i into a sorted list of indexes I. Get the predecessor and successor of u_i in I, called u_j and u_k respectively, if they exist (in which case $j < i$ and $k < i$). Use range search to check if u_j and u_i are contiguous, and if u_i and u_k are contiguous; if a pair does not exist, treat is as not being contiguous. If neither of these two pairs is contiguous, do nothing. If exactly one of these pairs is contiguous, use binary search in V to obtain the set of columns of E_1 embedded between the pair, add this set to B, and increment d. If both of these pairs are contiguous, use binary search to determine if there are elements of $\overline{E_1}$ (defined earlier as all the columns of class E_1) between u_j and u_i and between u_i and u_k; in which case we increment d. After we have iterated through $u(s_i)$ columns of U, it is time to add to $\hat{U}$ and create a new segment $\hat{s}_i \in \hat{S}$ with the following values.

$$u(\hat{s}_i) = |V| - |B|$$

$$d(\hat{s}_i) = d(s_i) + d + c_i$$

$$p(\hat{s}_i) = s_i$$

Remove all of the members of B from V, then add all of these members to the beginning of $\hat{U}$. Reset B to be the empty set. Continue iterating through U.

Once we have finished iterating through U, add the remaining elements in V to the start of $\hat{U}$.

Build the Next S-group. If E_2 is a row class, then $\hat{U} = \overline{E_1}$ and $\hat{S}$ contains one segment $\hat{s}$. To find $\hat{s}$ we iterate over each segment $s_i \in S$ looking for the segment s_i with minimum value $d(s_i) + c_i$. $\hat{s}$ has the following values:

$$u(\hat{s}) = |\overline{E_1}|$$

$$d(\hat{s}) = d(s_i) + c_i$$

$$p(\hat{s}) = s_i$$

If E_2 is a column class, $\hat{U}$ and $\hat{S}$ were created while we found embedded columns.

3.3 Merging Identical S-groups

Once we have generated a new S-group (E_1, E_2) we add it to a two dimensional table, where the row is determined by E_1's index in the original list of equivalence classes, and the column similarly determined from E_2's index.

If an S-group (E_1, E_2) already exists, we must merge the two S-groups. Given two (E_1, E_2) S-groups, $G_1 = \{U_1, S_1\}$ and $G_2 = \{U_2, S_2\}$, we will compute a new S-group, $G_3 = \{U_3, S_3\}$, that encompasses both of these S-groups which will then get stored in our table. We iterate through both G_1 and G_2 concurrently in order to create G_3, as described below. One starts with U_3 and S_3 being empty.

Merge S_1 and S_2, maintaining the ordering based on u Iterate through each s_i in this merged list. Let S_j be the list which contains s_i. Remove elements from the start of U_j, adding them to the end of U_3 until $|U_3| = u(s_i)$. (Do not add a value to U_3 if it is already in U_3.) This works because of the Containment Lemma from the Appendix **??**. Actually, the proof in [**ACJKLW07**] of the Containment Lemma relies on proving the fact that all the values of $u(s)$ with $s \in S_1$ are at most the minimum of the values of $u(s)$ with $s \in S_2$, or vice versa. Add s_i to S_3. If G_1 and G_2 both have segments such that $u(s_a) = u(s_b)$, choose the segment of smaller distance to add to S_3.

3.4 Finding the optimum SRRL

While we build the graph, we keep track of the segment which is a valid endpoint of the smallest distance. We define a valid endpoint to be a segment which is completely gray (this is since we solve the modified version of the problem; for the original problem, we would have a segment which is completely white and gray). A segment (E_1, E_2), is an endpoint if $E_1 = E_2$. (in the original version, we also have the case where E_1 and E_2 are adjacent in the equivalence class list and the cell where E_1 and E_2 intersect is white in the original pattern). Then once the graph is finished, we build a path to the valid endpoint using the values stored in p. This path will give the optimal order to remove equivalence classes, which can then be translated into the optimal list of rectangle rules.

This algorithms returns the optimum solution, as it follows from all the discussion above.

4 Time and Space Complexity

In this section we will show the time and space complexity of the previously defined algorithm is $O(nN \log n)$ and $O(nN)$ respectively – where N is the number of equivalence classes and n is the number of rows and columns in the original pattern. The ideas of our proof are partially taken from a more complete version of [**ACJKLW07**], which showed the number of reachable segments to be in $O(n^2)$. Indeed, our contribution is a faster way of processing a segment, cutting down this processing time down from $O(n)$ to $O(\log n)$.

Following [**ACJKLW07**], we will show that if the complexity for an S-group (E_a, E_b) is in $O(|E_a| + |E_b|)$, then the overall complexity of all S-groups is $O(nN)$. For each S-group, add the first term of its complexity, $O(|E_a|)$, to one two dimensional array, and its second, $O(|E_b|)$, to a second array – each at location (E_a, E_b).

$$\begin{bmatrix} E_1 & E_1 & \dots & E_1 \\ E_2 & E_2 & \dots & E_2 \\ \vdots & \vdots & \ddots & \vdots \\ E_N & E_N & \dots & E_N \end{bmatrix} \begin{bmatrix} E_1 & E_2 & \dots & E_N \\ E_1 & E_2 & \dots & E_N \\ \vdots & \vdots & \ddots & \vdots \\ E_1 & E_2 & \dots & E_N \end{bmatrix}$$

By noting that the sum of all members of all equivalence classes equals the total number of rows and columns, we get $\sum_{i=0}^{N} |E_i| = n$. The sum of the columns in the first matrix and the sum of the rows in the second matrix both equal n. The sum of all terms in both matrices is $2nN$, so the complexity of all S-groups is in $O(nN)$.

The space complexity of the algorithm is determined by the sizes of the lists storing the S-groups. Each S-group maintains a list U, which holds members of E_b, and is therefore in $O(|E_b|)$. The list S contains all segments. Each segment is guaranteed to have a unique size u, and the values of u are positive values less than $|E_b|$. The complexity of each segment is constant, so we again have $O(|E_b|)$. As we have previously shown, since each segment is in $O(|E_a| + |E_b|)$, the overall complexity is $O(nN)$.

Each operation we perform on an S-group (E_a, E_b) happens in either $O(|E_a| \log n)$ or $O(|E_b| \log n)$, as indeed every "check" from the algorithm's description takes time $O(\log n)$, after $O(n \log n$ initialization of the range search data structure or the ordered list represented by a balanced binary tree. So the runtime for processing the equivalence class list is $O(nN \log n)$.

Since reading the input takes $O(n^2)$ time and N is in $O(n)$ (this follows immediately from the Monotonicity Property, Theorem **??** in the Appendix), we relax our bounds and say our space complexity is $O(n^2)$ and the overall runtime is $O(n^2 \log n)$.

5 Conclusions

As noted in [**NLT13**; **NLT18**], these solutions do not generalize to dimensions higher than two. This is the most interesting open question, in our opinion.

We believe that range trees can be replaced by ad-hoc methods to obtain a $O(n^2)$ algorithm for exact RRL strip-rule minimization. The savings of $O(\log n)$ in running time comes at the expense of a more complicated algorithm which we decided not to present.

Acknowledgements

6 Appendix

6.1 Correctness of Pick-Up-Sticks Algorithm

For the sake of completeness, we include:

Theorem 1 (part of Theorem 3.1 of [ACJKLW07]). *A black and white pattern P is a strip-rule pattern if and only if the Pick-Up-Sticks algorithm results in a grid with all of its cells gray. In that case, the reverse list of picked-up rows and columns is a SRRL for P.*

Since Theorem 3.1 of [**ACJKLW07**] is not proven in their conference version, we include a proof for the sake of completeness.

Proof. Let's separate the proof into the two following parts:

- (P is a strip-rule pattern) $\rightarrow$ (The Pick-Up-Sticks algorithm results in an all-gray grid)
  Suppose P is a pattern that can be generated by a SRRL $S = (s_1, s_2, \cdots, s_m)$ of size m. Suppose we have picked up some sticks. Let s_i is the highest-index rule from S that still has a row/column c such that c is not all-gray. If there is no such s_i, then we must have picked up everything. For every integer j, $i < j \leq m$, all the rows/columns of s_j must have already been entirely colored gray. That means that all cells of c that belong to a rule s_j, with $i < j \leq m$, are gray. The cells of c that are not gray must have the same color as the rule s_i, as they will not be overwritten by any rule s_j with $i < j \leq m$. Hence, c must be pseudo-monochromatic and then the Pick-Up-Sticks may proceed. As a result, the algorithm never gets stuck until the all-gray grid is reached.
- (The Pick-Up-Sticks algorithm results in an all-gray grid) $\rightarrow$ (P is a strip-rule pattern)
  The reverse list of the picked-up rows and columns is a SRRL since on every iteration of the algorithm we only pick up strips of P. If the Pick-Up-Sticks algorithm results in an all-gray grid, every row and column of P has been picked up by the end of the algorithm. Hence, the reversed list of picked up rows and columns is a SRRL that generates P, making P a strip-rule pattern.

  This ends our proof. $\square$

The same proof works (the number of rules may differ by one) if the Pick-Up-Sticks algorithm results in a grid with none of its cells black.

6.2 The Monotonicity Property

Theorem 2 (part of Theorem 3.1 of [ACJKLW07]). *(Monotonicity Property) Let P be a black and white pattern. P is a strip-rule pattern if and only if the two equivalent properties holds:*

(i) *For any color $c \in \{black, white\}$, the rows of P are hierarchical: given any two rows of P, the set of columns where c is present in the first row either contains or is contained in the set of columns where c is present in the second.*
(ii) *The same property of item (i) holds with the roles of row and column switched.*

**Fig. 7.** A pattern that follows the Monotonicity Property

Fig. 8. A pattern that does not follow the Monotonicity Property

This theorem will allow us to order the columns and rows of a strip-rule pattern P in a non-decreasing sequence regarding the number of black cells. Once the i-th black cell is achieved in the sequence, that i-th cell will remain black until the end of the sequence.

Since Theorem 3.1 of [**ACJKLW07**] is not proven in their conference version, we include a proof for the sake of completeness.

Proof. Let's begin our proof by verifying that items (i) and (ii) of the Monotonicity Property are equivalent.

I.) $(i) \leftrightarrow (ii)$

Let's start by proving that $\neg(i) \to \neg(ii)$. Suppose (i) does not hold. Then, given $c \in \{black, white\}$, there are two rows A and B , $A \neq B$, of P that do no follow the property mentioned above. Without loss of generality, let's say $c = black$. Since one set is not contained into each other, it must be the case where both sets contains a element that the other set does not have.

This implies that both of the following are true at the same time:

1. There is a column that has a black cell on row A and a white cell on row B.
2. There is a column that has a black cell on row B and a white cell on row A.

For reference, let's call this case a checkerboard case.

Fig. 9. Example of a checkerboard case, where (i) does not hold.

Let X and Y be one of those columns described in (a) and (b) respectively. The set of rows where black is present in X has A as one of its rows and the set of rows where black is present in Y has B as one of its rows. Hence, those sets are not contained into each other. As a result, (ii) does not hold. Thus it must be true that (ii) implies (i). $((ii) \to (i))$

The same argument can be made with the roles of rows and columns switched, leading to a conclusion that (i) implies (ii). $((i) \to (ii))$ Then it must be true that (i) and (ii) are equivalent.

Let's continue with the equivalence of P being a strip-rule pattern and the Monotonicity property. Let's divide the proof into two steps:

– (P is a strip-rule pattern) $\to$ (Monotonicity Property)

We will prove that ¬(Monotonicity Property) → ¬(P is a strip-rule pattern). Suppose the Monotonicity property does not hold true. By contradiction, suppose that P is a strip-rule pattern and let S be a strip-rule RRL of size n that generates P such that $S = (s_1, s_2, \cdots, s_n)$.

Since the Monotonicity property does not hold true, the same argument as before for the checkerboard case can be made. Let s_i be the last rule of S that includes one of the rows or columns of the checkerboard case. Since s_i is the last rule that they appear, no other rule after s_i could have changed the intersection of the rows A, B and columns X, Y of the checkerboard case. Then, because of s_i, either A, B, X or Y must have same color cells in the checkerboard. However, if that happens, the checkerboard case does not hold anymore - a contradiction. Hence, P is not a strip-rule pattern and (P is a strip-rule pattern) → (Monotonicity Property) must hold true.

– (Monotonicity Property) → (P is a strip-rule pattern)

Suppose that the Monotonicity property holds. By contradiction, suppose that P is not a strip-rule pattern. By **??**, the MPUS algorithm will not stop with a all-gray grid. Then it must have been stuck at a grid with no pseudo-monochromatic row or column.

Let A be the row with the largest number of cells colored black. Since this row is not monochromatic, there must be a column Y which intersection with A is a white cell. Y must also not be pseudo-monochromatic. So there must a row B that intersects Y with a black cell. Since B has a black cell in a column that A does not have, than it must be true that the black columns of A are a subset of the black columns of B because of the Monotonicity property. Note also that those sets must not be equal.

Hence, the number of black cells in B is higher than A, a contradiction. As a conclusion, (Monotonicity Property) → (P is a strip-rule pattern) must hold, which concludes our proof.

In conclusion, have shown that patterns generated by RRLs and patterns that are satisfy monotonicity are equivalent. This allows us to label the row and column classes by the number of black cells they contain. For example, every row that has exactly 4 black cells will be identical and can be grouped into the class named R_4. Our algorithm uses this principle to organize and determine the order of classes to be removed. The fact that RRL patterns have all the properties of monotonicity will helpful is number of following proofs. □

6.3 Proofs Regarding Equivalence Classes

This section contains some of the proofs relevant to the equivalence classes ordering and the creation of the hierarchical array of a pattern. We point out which of these appear in [**ACJKLW07**], and include proofs for completeness' sake, since [**ACJKLW07**] omits these proofs.

Active Classes

Proposition 1 ([ACJKLW07], Observation 6. on page 1070). *During the execution of the MPUS algorithm on a black and white strip-rule pattern, there are always exactly two active classes at any given time. The two active classes are either a row and a column class of the same color or both classes of the same kind (rows or columns), being one of each color.*

Since Observation 6 of [**ACJKLW07**] is not proven in their conference version, we include a proof for the sake of completeness.

Proof. First, we argue that at least two classes are active. We use the knowledge that the Pick-Up-Sticks (PUS) algorithm will always terminate on a gray board if an RRL solution exists. Suppose at any time during PUS, there is only one class active. Without loss of generality, assume this active class is a set of black rows. We will consider the impact of picking up a black row on other potential classes becoming active. Picking up a black row has no impact on whether a white row becomes active. It has no impact on a black column becoming available – this class needs white rows to be removed. Finally, only once all of the black rows have been removed can a white column class become active. This means that picking up some or all of the active class can only result in there either being zero active classes or still only one being active. Contrast this with the final state of the PUS algorithm, a solid board which translates to two active classes (a row and column of that color). Therefore, a board with only one active class can never get to the state with two active classes required to terminate PUS and be solvable by RRL.

Next, we argue that the two active classes are either a row and a column class of the same color or both classes of the same kind (rows or columns), being one of each color. Indeed, let us first look at the case when the two active classes are both column classes (with a symmetric case with two row classes), and let there be two columns D_1 and D_2 one each from these two active classes, such that D_1 and D_2 have not been picked up. There is a row

where D_1 has a black cell and D_2 has a white cell; these cells cannot be gray since neither D_1 nor D_2 have been picked up, and the row has never been pseudo-monochromatic. This means D_1 can only be pseudo-monochromatic as a black-gray column, and D_2 can only be pseudo-monochromatic as a white-gray column. Second, let us look at the case where the two active classes are one a row class and a column class, and let D_1 be a row from the first active class and let D_2 be a column from the second active class, such that D_1 and D_2 have not been picked up. Then the cell at the intersection of D_1 and D_2 is not gray, and so both D_1 and D_2 are pseudo-monochromatic with the color of this cell.

Note that the fact proved in the previous paragraph implies that there cannot be three or more active classes: there are not enough colors for three active column classes, while two active columns classes and an active two class will lead to a contradiction since both columns classes are of the same color as the row class, while being at the same time of different colors. Thus exactly two classes are active at any given time. □

Proposition 2 (Implicit in [ACJKLW07]). *For a pattern P of two colors, if one picks up all columns of an active class, a row class of the opposite color becomes active. If one picks up all rows of an active class, a column class of the opposite color becomes active.*

Proof. Since P is a pattern composed of only two colors, we have that there are always two active classes at a time. Picking up a (monochromatic) column active class will not change the monochromatic property of any other column class. Hence, the class that will become monochromatic needs to be a row class. The same thing holds with the roles of rows and columns switched.

Also, the new active class must be of the opposite color because if it was not, the class becoming active would need to be already active before, since the color being removed is the same as the other cells of that class, which means that it would need to be monochromatic already. □

Number of Equivalence Classes

Proposition 3 (Implicit in [ACJKLW07]). *Let i be such that $1 \leq i < N_C$ and j be such that $1 \leq j < N_R$.*

a) *Let S_i be the set of rows that have white cells on the intersections of the columns of C_{x_i} and black cells on intersection of the columns of $C_{x_{i+1}}$. S is a row class ($S_i \in R$).*
b) *Let S_j be the set of columns that have white cells on the intersections of the rows of R_{y_j} and black cells on intersection of the columns of $R_{y_{j+1}}$. S_j is a column class ($S_j \in C$).*

Proof (Proof of item (a)). Let i be such that $1 \leq i < N_C$ and S be as the set described in the statement of the proposition. Let $\alpha, \beta \in S$. By contradiction, suppose α and β differ from each other. Let γ be one of the columns that they differ and $k \in \mathbb{N}^*$ be such that C_{x_k} is the class of γ. Without loss of generality, assume C_{x_k} has a black cell on the row α and a white cell on the row β.

Because of monotonicity, the fact that C_{x_k} has a black cell on α and C_{x_i} has a white cell on α implies that $x_k > x_i$ must hold true. At the same time, the fact that C_{x_k} has a white cell on β and $C_{x_{i+1}}$ has a black cell on β implies that $x_k < x_{i+1}$ must hold true. Hence, $x_i < x_k < x_{i+1}$.

However, $x_i < x_k < x_{i+1}$ implies that $i < k < i+1$, which is a contradiction because $k \in \mathbb{N}^*$. ⊥

Hence, α and β do not differ from each other, which implies that S_i is a row class ($S_i \in R$).

The proof of item (b) follows analogously with the roles of rows and columns switched. □

Proposition 4. *The number of row classes and the number of column classes differ by one. That is $|N_R - N_C| \leq 1$.*

Proof. Let $A = max(N_R, N_C)$ and $B = min(N_R, N_C)$. Note that $|N_R - N_C| = A - B$. If $N_R = N_C$, we have that $|N_R - N_C| = |0| = 0 \leq 1$.

Now, suppose $N_R \neq N_C$. Assume $A = N_C$. It follows that $B = N_R$.

Let i_1 be such that $1 \leq i_1 < N_C$. Let S_{i_1} be the set of rows that have white cells on the intersections of the columns of $C_{x_{i_1}}$ and black cells on intersection of the columns of $C_{x_{i_1+1}}$. According to **??** we have that S_{i_1} is a row class.

Let i_2 be such that $1 \leq i_2 < N_C$ and $i_2 \neq i_1$. Let S_{i_2} be the set of rows that have white cells on the intersections of the columns of $C_{x_{i_2}}$ and black cells on intersection of the columns of $C_{x_{i_2+1}}$. According to **??**(a) we have that S_{i_2} is a row class.

If $i_2 < i_1$, we have that $x_{i_2+1} \leq x_{i_1}$, which means that the rows of S_{i_2} must have a black cell on its intersections with $C_{x_{i_1}}$ due to monotonicity. This implies that $S_{i_2} \neq S_{i_1}$ because they differ on $C_{x_{i_1}}$.

If $i_2 > i_1$, we have that $x_{i_2} \geq x_{i_1+1}$, which means that the rows of S_{i_2} must have a white cell on its intersections with $C_{x_{i_1+1}}$ due to monotonicity. This implies that $S_{i_2} \neq S_{i_1}$ because they differ on $C_{x_{i_1+1}}$.

Thus, for every i such that $1 \leq i < N_C$ we have a different S_i that is a row class. From that, we have that $N_R \geq N_C - 1$, being $N_C - 1$ the number of different possibilities for i. Hence, $N_C - N_R = A - B \leq 1$.

Now assume $A = N_R$. It follows that $B = N_C$. Using the same reasoning with **??**(b), we have that $N_R \geq N_C - 1$. Hence, $N_R - N_C = A - B \leq 1$.

Therefore, it is true for any case that $|N_R - N_C| \leq 1$. □

Ordering

Proposition 5. *There is an array E_1, E_2, $\cdots$, E_N of equivalence classes that respect the following property: given a and b such that $1 \leq a < b \leq N$, if E_a and E_b are the active classes, then for every c such that $1 \leq c < a$ or $b < c \leq N$, E_c has already been picked up. Also, if $b = a+1$, then either $E_a \in R$ and $E_b \in C$ or $E_a \in C$ and $E_b \in R$. Moreover, none of the columns/rows of a class E_c for $a < c < b$ are pseudo-monochromatic. This array is the hierarchical array of a pattern.*

Proof. We know from **??** that there are always exactly 2 acti ve classes during an execution of PUS. We combine this idea with the concept of monotonicity in order to create an organization on the possible orders the classes are picked up in. For the original pattern, let A and B be the two original active classes . If $A, B \in C$, we must have that one is white and the other is black. Assume without loss of generality that A is white and B is black. Then we have that $A = C_{x_1} = C_0$ and $B = C_{x_{N_C}} = C_{n_R}$. To build the array, we will assign $E_1 = A = C_{x_1}$ and $E_N = B = C_{x_{N_C}}$ since A and B need to be on both ends of the array because bo th may be active even if no other classes have been picked. If we pick up C_{x_1}, a row class will need to become active. Also, this class will need to be the row class with the largest number of bl ack cells due to monotonicity. Hence, we should assign $E_2 = R_{y_{N_R}}$. Following the same reasoning, the next member will need to be the column cla ss with the lowest number of black cells after C_{x_1}. Hence, we should assign $E_3 = C_{x_2}$. Analogously, if we pick up $C_{x_{N_C}}$, the row class with the least amo unt of black cells will become available, this being R_{y_1}. Hence, we should assign $E_{N-1} = R_{y_1}$, $E_{N-2} = C_{x_{N_C-1}}$ a nd so on.

Therefore we can build this array assigning $E_{2i-1} = C_{x_i}$ and $E_{2j} = R_{y_{N_R-j+1}}$ for every $1 \leq i \leq N_C$ and $1 \leq j \leq N_R$ with $E_N = C_{x_{N_C}}$. Given a and b such that $1 \leq a < b \leq N$, we have that if E_a a nd E_b are the active classes, every class that comes before E_a or after E_b must have already been picked up. Also, note that two consecutive elements of the array are not of the same ty pe of equivalence classes (row/column). Note also that that $E_{2N_C-1} = C_{x_{N_C}} = E_N$, which means th at $N = N_R + N_C = 2N_C - 1$ and $N_R + 1 = N_C$, which agrees w ith **??**. If $A, B \in R$ we may build this array analogously assigning $E_{2j-1} = R_{y_j}$ and $E_{2i} = C_{x_{N_C-i+1}}$ for every $1 \leq i \leq N_C$ a nd $1 \leq j \leq N_R$ and $E_N = R_{y_{N_R}}$. For this case, $N_C + 1 = N_R$.

For A and B being of different types, let's assume the color of the activ e classes, which is the same, is white. We can build this array assigning $E_{2i-1} = C_{x_i}$ and $E_{2j} = R_{y_{N_R-j+1}}$ for every $1 \leq i \leq N_C$ and $1 \leq j \leq N_R$ wi th $E_N = R_{y_1}$. For this case, $E_{2N_R} = R_{y_1} = E_N$ and $N = N_R + N_C = 2N_R$ and $N_R = N_C$, which agrees with **??**. If the color is black, we can build this array assigning $E_{2j-1} = R_{y_j}$ and $E_{2i} = C_{x_{N_C-i+1}}$ for every $1 \leq i \leq N_C$ and $1 \leq j \leq N_R$ with $E_N = C_{x_{N_C}}$. For this case, $E_{2N_C} = C_{x_{N_C}} = E_N$ and $N = N_R + N_C = 2N_C$ and $N_R = N_C$. By construction it is easy to see that a column/row of a class E_c must no t be monochromatic while the two active classes are E_a and E_b for $a < c < b$. From our construction and definitions, the only way for a class to become ac tive, is to have either the class immediately to its left or right in the hi erarchical array already removed. This therefore cannot be true while $a < c < b$ since we know the only class es removed are those that come before a or come after b. □

6.4 Segments Properties

This section contains the proofs of the properties discussed earlier regarding segments.

Segment Redefinition Earlier, we mentioned that the definition given by "Applegate et al."'s paper [**ACJKLW07**] is redundant. They define a segment as being a tuple of five elements (E_x, c_x, E_y, c_y, U), where E_x and E_y are the active equivalence classes, c_x and c_y are the colors of the members E_x and E_y (that are pseudo-monochromatic), and U is the subset of members left unpicked of E_y ($U \subset E_y$). The next statement shows that the colors are redundant, allowing us to reduce the definition of a segment to $(E_x, E_y, , U)$.

Proposition 6. *Let E_1, E_2, $\cdots$, E_N be the hierarchical array of a pattern. Given a and b such that $1 \leq a < b \leq N$ it is possible to know what colors correspond to each active class of a label (E_a, E_b).*

Proof. Given a and b such that $1 \leq a < b \leq N$.

- If $E_1 \in C$ and $E_N \in C$, we have that: If $E_a \in C$, E_a is white. If $E_a \in R$, E_a is black. If $E_b \in C$, E_b is black. If $E_b \in R$, E_a is white.
- If $E_1 \in R$ and $E_N \in R$, we have that: If $E_a \in C$, E_a is black. If $E_a \in R$, E_a is white. If $E_b \in C$, E_b is white. If $E_b \in R$, E_a is black.
- If E_1 and E_N are from different types and E_1 and E_N are initially black, we have that: If $E_a \in C$, E_a is black. If $E_a \in R$, E_a is white. If $E_b \in C$, E_b is white. If $E_b \in R$, E_a is black.
- If E_1 and E_N are from different types and E_1 and E_N are initially white, we have that: If $E_a \in C$, E_a is white. If $E_a \in R$, E_a is black. If $E_b \in C$, E_b is black. If $E_b \in R$, E_a is white.

This happens because when we pick from either sides, the color alternates for each item picked up as well as its type. Hence, the color of E_a will always match the color of E_1 if they have the same type and differ if they are from different types. The same will happen with E_b and E_N. □

It is a consequence of this corollary the fact that having a segment $S = (E_{k_1}, c_{k_1}, E_{k_2}, c_{k_2}, U)$, $U \subset E_{k_2}$, is redundant. We may start referring to a segment as only $S = (E_{k_1}, E_{k_2}, U)$, $U \subset E_{k_2}$.

Proposition 7 (statement 2(a) of Lemma 4.2 of [ACJKLW07]). *Suppose we have two sets U_x and U_y of active blocks of the two active classes E_x and E_y. Suppose there is an optimal solution that completely picks up E_x before completely picking up E_y. Then there also exists an optimal solution that picks up all embedded members of E_y before picking up E_x.*

Since Lemma 4.2 of [**ACJKLW07**] is not proven in their conference version, we include a proof for the sake of completeness.

Proof. This follows closely from the definition of embedded members, which can be effectively picked up at zero net cost along with the blocks embedding them. This is easy to observe in an example: consider the leftmost three columns of Fig. **??**. Here we see a member of C_0 (the white column) embedded in two blocks of C_2. We could pick up the embedding blocks using two rules, or we could first pick up the embedded block followed by picking up both embedding blocks using only one wider rule. The overall result is still two rules which pick up all members of the embedding class. This is, without loss of generallity, either a helpful or irrelevant action. So in our solution we will always choose to remove the embedded blocks before the embedding blocks. □

Containment Lemma [Lemma 4.3 of [**ACJKLW07**]] Let S and S' be two segments on the same equivalence classes, then either $U \subset U'$ or $U' \subset U$.

Intuitively, this follows from the fact that we have a strict ordering on the equivalence classes that could embed columns in U. See "Applegate et al."'s paper [**ACJKLW07**] for a rigorous inductive proof of this lemma by the same name.